

```
In [1]:  import pandas as pd  # data frame operations
import numpy as np  # arrays and math functions
import matplotlib.pyplot as plt  # static plotting
```

```
In [28]:  # correlation heat map setup for seaborn
def corr_chart(df_corr):
    corr=df_corr.corr()
    #screen top half to get a triangle
    top = np.zeros_like(corr, dtype=np.bool)
    top[np.triu_indices_from(top)] = True
    fig=plt.figure()
    fig, ax = plt.subplots(figsize=(12,12))
    sns.heatmap(corr, mask=top, cmap='coolwarm',
                center = 0, square=True,
                linewidths=.5, cbar_kws={'shrink':.5},
                annot = True, annot_kws={'size': 9}, fmt = '.3f')
    plt.xticks(rotation=45) # rotate variable labels on columns (x axis)
    plt.yticks(rotation=0) # use horizontal variable labels on rows (y axis)
    plt.title('Correlation Heat Map')
    plt.savefig('plot-corr-map.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)
```

```
In [5]:  valid_survey_input = pd.read_csv('mspa-survey-data.csv')

# use the RespondentID as label for the rows... the index of DataFrame
```

```
In [7]:  print('\nContents of initial survey data -----')
```

Contents of initial survey data -----

Number of Respondents = 207

```
In [8]: print(valid_survey_input.columns)
```

```
Index(['Personal_JavaScalaSpark', 'Personal_JavaScriptHTMLCSS',  
      'Personal_Python', 'Personal_R', 'Personal_SAS',  
      'Professional_JavaScalaSpark', 'Professional_JavaScriptHTMLCSS',  
      'Professional_Python', 'Professional_R', 'Professional_SAS',  
      'Industry_JavaScalaSpark', 'Industry_JavaScriptHTMLCSS',  
      'Industry_Python', 'Industry_R', 'Industry_SAS',  
      'Python_Course_Interest', 'Foundations_DE_Course_Interest',  
      'Analytics_App_Course_Interest', 'Systems_Analysis_Course_Interest',  
      'Courses_Completed', 'PREDICT400', 'PREDICT401', 'PREDICT410',  
      'PREDICT411', 'PREDICT413', 'PREDICT420', 'PREDICT422', 'PREDICT450',  
      'PREDICT451', 'PREDICT452', 'PREDICT453', 'PREDICT454', 'PREDICT455',  
      'PREDICT456', 'PREDICT457', 'OtherPython', 'OtherR', 'OtherSAS',  
      'Other', 'Graduate_Date'],  
      dtype='object')
```

```
In [9]: print(pd.DataFrame.head(valid_survey_input))
#Here we can see the first few values of each of the columns. This illustrates how students distribute the poi.
```

RespondentID	Personal_JavaScalaSpark	Personal_JavaScriptHTMLCSS	\
5135740122	0	0	
5133300037	10	10	
5132253300	20	0	
5132096630	10	10	
5131990362	20	0	

RespondentID	Personal_Python	Personal_R	Personal_SAS	\
5135740122	0	50	50	
5133300037	50	30	0	
5132253300	40	40	0	
5132096630	25	35	20	
5131990362	0	70	10	

RespondentID	Professional_JavaScalaSpark	Professional_JavaScriptHTMLCSS	\
5135740122	0	0	
5133300037	25	25	
5132253300	0	0	
5132096630	10	10	
5131990362	20	0	

RespondentID	Professional_Python	Professional_R	Professional_SAS	\
5135740122	0	25	75	
5133300037	30	20	0	
5132253300	40	40	20	
5132096630	25	35	20	
5131990362	0	80	0	

RespondentID	...	PREDICT453	PREDICT454	PREDICT455	PREDICT456	\
5135740122	...	NaN	NaN	NaN	NaN	
5133300037	...	NaN	NaN	NaN	NaN	
5132253300	...	NaN	NaN	NaN	NaN	
5132096630	...	NaN	NaN	NaN	NaN	
5131990362	...	NaN	NaN	NaN	NaN	

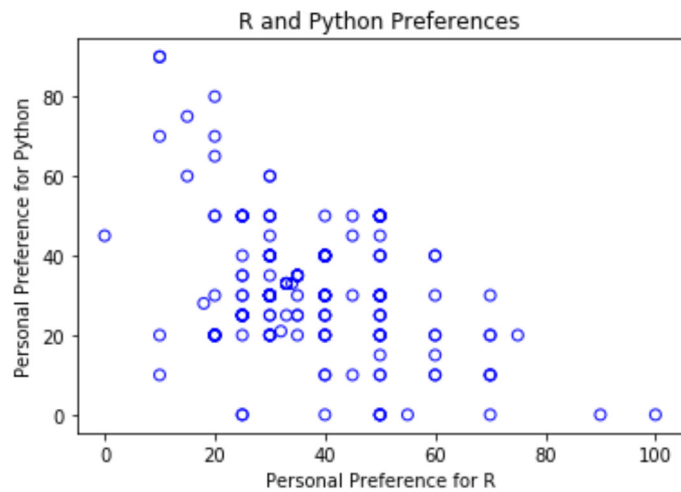
RespondentID	PREDICT457	OtherPython	OtherR	OtherSAS	Other	\
5135740122	NaN	NaN	NaN	NaN	NaN	
5133300037	NaN	NaN	NaN	NaN	NaN	

```
In [10]: survey_df = valid_survey_input.rename(index=str, columns={
    'Personal_JavaScalaSpark': 'My_Java',
    'Personal_JavaScriptHTMLCSS': 'My_JS',
    'Personal_Python': 'My_Python',
    'Personal_R': 'My_R',
    'Personal_SAS': 'My_SAS',
    'Professional_JavaScalaSpark': 'Prof_Java',
    'Professional_JavaScriptHTMLCSS': 'Prof_JS',
    'Professional_Python': 'Prof_Python',
    'Professional_R': 'Prof_R',
    'Professional_SAS': 'Prof_SAS',
    'Industry_JavaScalaSpark': 'Ind_Java',
    'Industry_JavaScriptHTMLCSS': 'Ind_JS',
    'Industry_Python': 'Ind_Python',
    'Industry_R': 'Ind_R',
```

```
In [11]:
```

```
In [13]: fig, axis = plt.subplots()
axis.set_xlabel('Personal Preference for R')
axis.set_ylabel('Personal Preference for Python')
plt.title('R and Python Preferences')
scatter_plot = axis.scatter(survey_df['My_R'],
                             survey_df['My_Python'],
                             facecolors = 'none',
                             edgecolors = 'blue')
plt.savefig('plot-scatter-r-python.pdf',
            bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
            orientation='portrait', papertype=None, format=None,
            transparent=True, pad_inches=0.25, frameon=None)

#This graph shows us that the majority of the values in the personal category fall in the center of this graph
#The center of this graph contains numbers between 20-50 for python and 20-60 for R.
```




In [14]:

```

survey_df_labels = [
    'Personal Preference for Java/Scala/Spark',
    'Personal Preference for Java/Script/HTML/CSS',
    'Personal Preference for Python',
    'Personal Preference for R',
    'Personal Preference for SAS',
    'Professional Java/Scala/Spark',
    'Professional JavaScript/HTML/CSS',
    'Professional Python',
    'Professional R',
    'Professional SAS',
    'Industry Java/Scala/Spark',
    'Industry Java/Script/HTML/CSS',
    'Industry Python',
    'Industry R',
    'Industry SAS'
]


# create a set of scatter plots for personal preferences
for i in range(5):
    for j in range(5):
        if i != j:
            file_title = survey_df.columns[i] + '_and_' + survey_df.columns[j]
            plot_title = survey_df.columns[i] + ' and ' + survey_df.columns[j]
            fig, axis = plt.subplots()
            axis.set_xlabel(survey_df_labels[i])
            axis.set_ylabel(survey_df_labels[j])
            plt.title(plot_title)
            scatter_plot = axis.scatter(survey_df[survey_df.columns[i]],
                                        survey_df[survey_df.columns[j]],
                                        facecolors = 'none',
                                        edgecolors = 'blue')
            plt.savefig(file_title + '.pdf',
                        bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                        orientation='portrait', papertype=None, format=None,

```



```
In [15]: corr_chart(df_corr = software_df)
#There are too many relationships between variables here to analyze briefly. But we can say that the
#stronger correlations in the graph belong to relationships of the same languages. This means that
#respondents' desire to learn languages(My_....) is strongly correlated with how important the languages are
#for the jobs they are aspiring to obtain and also it's strongly correlated with the languages that are
#relevant in the industry.
#The rest of the correlations are not significant in my opinion since I consider minimum 0.65 to
#be a good correlation but the closest one that doesn't fall in the category previously described is the
#relationship between Industry Python and Industry SAS, the number tells us that about 60% of respondents
#that showed high interest in Python also showed in SAS.
```

<Figure size 432x288 with 0 Axes>



```
In [16]: ▶ print('\nDescriptive statistics for survey data -----')
print(software_df.describe())
#The descriptive stats below give Python and R a big win with R in first place and SAS comes behind these two.
#The values for relevance for languages across the 3 questions vary slightly which tells us that
#the desire for respondents is almost perfectly linearly related to what the industry and the jobs
```

Descriptive statistics for survey data -----

	My_Java	My_JS	My_Python	My_R	My_SAS	Prof_Java \
count	207.000000	207.000000	207.000000	207.000000	207.000000	207.000000
mean	10.135266	4.797101	31.304348	37.125604	16.637681	9.251208
std	11.383477	6.757764	15.570982	14.576003	13.626400	13.167505
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	20.000000	30.000000	5.000000	0.000000
50%	9.000000	0.000000	30.000000	35.000000	15.000000	5.000000
75%	20.000000	10.000000	40.000000	50.000000	25.000000	15.000000
max	70.000000	30.000000	90.000000	100.000000	75.000000	80.000000

	Prof_JS	Prof_Python	Prof_R	Prof_SAS	Ind_Java \
count	207.000000	207.000000	207.000000	207.000000	207.000000
mean	5.840580	30.028986	36.415459	18.463768	11.942029
std	10.812555	19.144802	20.847606	18.831841	14.706399
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	20.000000	25.000000	0.000000	0.000000
50%	0.000000	30.000000	33.000000	15.000000	5.000000
75%	10.000000	40.000000	50.000000	30.000000	20.000000
max	100.000000	100.000000	100.000000	100.000000	70.000000

	Ind_JS	Ind_Python	Ind_R	Ind_SAS
count	207.000000	207.000000	207.000000	207.000000
mean	6.966184	29.772947	32.434783	18.884058
std	10.030721	17.959816	15.912209	19.137623
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	20.000000	22.500000	0.000000
50%	0.000000	30.000000	30.000000	15.000000
75%	10.000000	40.000000	40.000000	30.000000
max	50.000000	95.000000	85.000000	100.000000


```
In [32]: ▶ classes_df = survey_df.loc[:, 'Python_Course_Interest':'Systems_Analysis_Course_Interest']

print('\nDescriptive statistics for survey data -----')
print(classes_df.describe())
```

```
Descriptive statistics for survey data -----
      Python_Course_Interest  Foundations_DE_Course_Interest  \
count                206.000000                      200.000000
mean                 73.529126                      58.045000
std                  29.835429                      32.588079
min                   0.000000                      0.000000
25%                  53.000000                      29.500000
50%                  82.500000                      60.000000
75%                 100.000000                      89.250000
max                 100.000000                      100.000000

      Analytics_App_Course_Interest  Systems_Analysis_Course_Interest
count                203.000000                      200.000000
mean                 55.201970                      53.630000
std                  34.147954                      33.539493
min                   0.000000                      0.000000
25%                  25.000000                      21.500000
50%                  60.000000                      51.500000
75%                  85.000000                      80.250000
max                 100.000000                      100.000000
```

```
In [17]: ▶ print('\nDescriptive statistics for courses completed -----')
print(survey_df['Courses_Completed'].describe())
```

*#These stats tell us that 20 people did not complete this part of the survey because there were 207
#responses as shown above and only 187 are shown on the count below. It also tell us that the average*

```
Descriptive statistics for courses completed -----
count    187.000000
mean      6.342246
std       3.170849
min       1.000000
25%       4.000000
50%       6.000000
75%       9.000000
max      12.000000
Name: Courses_Completed, dtype: float64
```

```
In [18]: from sklearn.preprocessing import StandardScaler
```

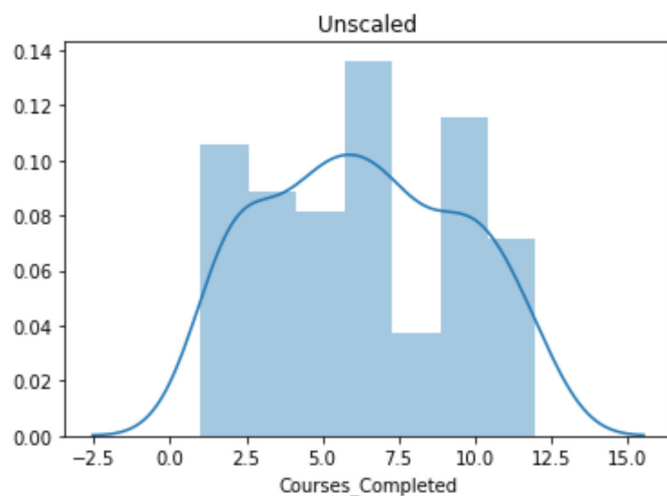
```
In [20]: X = survey_df['Courses_Completed'].dropna()  
len(X)
```

```
Out[20]: 187
```

```
In [24]: unscaled_fig, ax = plt.subplots()  
sns.distplot(X).set_title('Unscaled')  
unscaled_fig.savefig('Transformation-Unscaled' + '.pdf',  
                    bbox_inches='tight', dpi=None, facecolor='w', edgecolor='b',  
                    orientation='portrait', papertype=None, format=None,  
                    transparent=True, pad_inches=0.25, frameon=None)
```

C:\Users\gilad\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



```
In [25]:  standard_fig, ax = plt.subplots()
sns.distplot(StandardScaler().fit_transform(X)).set_title('StandardScaler')
standard_fig.savefig('Transformation-StandardScaler' + '.pdf',
                    bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                    orientation='portrait', papertype=None, format=None,
```

ValueError

Traceback (most recent call last)

<ipython-input-25-2740750c98f4> in <module>

```
1 standard_fig, ax = plt.subplots()
----> 2 sns.distplot(StandardScaler().fit_transform(X)).set_title('StandardScaler')
3 standard_fig.savefig('Transformation-StandardScaler' + '.pdf',
4     bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
5     orientation='portrait', papertype=None, format=None,
```

~\Anaconda3\lib\site-packages\sklearn\base.py in fit_transform(self, X, y, **fit_params)

```
460     if y is None:
461         # fit method of arity 1 (unsupervised transformation)
--> 462         return self.fit(X, **fit_params).transform(X)
463     else:
464         # fit method of arity 2 (supervised transformation)
```

~\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py in fit(self, X, y)

```
623     # Reset internal state before fitting
624     self._reset()
--> 625     return self.partial_fit(X, y)
626
627     def partial_fit(self, X, y=None):
```

~\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py in partial_fit(self, X, y)

```
647     X = check_array(X, accept_sparse=('csr', 'csc'), copy=self.copy,
648                     warn_on_dtype=True, estimator=self, dtype=FLOAT_DTYPES,
--> 649                     force_all_finite='allow-nan')
650
651     # Even in the case of `with_mean=False`, we update the mean anyway
```

~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_array(array, accept_sparse, accept_large_s
parse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, wa
rn_on_dtype, estimator)

```
550     "Reshape your data either using array.reshape(-1, 1) if "
551     "your data has a single feature or array.reshape(1, -1) "
--> 552     "if it contains a single sample.".format(array))
553
554     # in the future np.flexible dtypes will be handled like object dtypes
```

ValueError: Expected 2D array, got 1D array instead:

array=[6. 4. 7. 7. 5. 11. 2. 3. 6. 3. 2. 7. 3. 4. 2. 12. 7. 5.]

```
In [26]: minmax_fig, ax = plt.subplots()
sns.distplot(MinMaxScaler().fit_transform(X)).set_title('MinMaxScaler')
minmax_fig.savefig('Transformation-MinMaxScaler' + '.pdf',
                  bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                  orientation='portrait', papertype=None, format=None,
```

ValueError

Traceback (most recent call last)

<ipython-input-26-ea06a6d08a5f> in <module>

```
1 minmax_fig, ax = plt.subplots()
----> 2 sns.distplot(MinMaxScaler().fit_transform(X)).set_title('MinMaxScaler')
3 minmax_fig.savefig('Transformation-MinMaxScaler' + '.pdf',
4     bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
5     orientation='portrait', papertype=None, format=None,
```

~\Anaconda3\lib\site-packages\sklearn\base.py in fit_transform(self, X, y, **fit_params)

```
460     if y is None:
461         # fit method of arity 1 (unsupervised transformation)
--> 462         return self.fit(X, **fit_params).transform(X)
463     else:
464         # fit method of arity 2 (supervised transformation)
```

~\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py in fit(self, X, y)

```
321     # Reset internal state before fitting
322     self._reset()
--> 323     return self.partial_fit(X, y)
324
325     def partial_fit(self, X, y=None):
```

~\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py in partial_fit(self, X, y)

```
349     X = check_array(X, copy=self.copy, warn_on_dtype=True,
350                     estimator=self, dtype=FLOAT_DTYPES,
--> 351                     force_all_finite="allow-nan")
352
353     data_min = np.nanmin(X, axis=0)
```

~\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_array(array, accept_sparse, accept_large_s
parse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_min_features, wa
rn_on_dtype, estimator)

```
550         "Reshape your data either using array.reshape(-1, 1) if "
551         "your data has a single feature or array.reshape(1, -1) "
--> 552         "if it contains a single sample.".format(array))
553
554     # in the future np.flexible dtypes will be handled like object dtypes
```


ValueError: Expected 2D array, got 1D array instead:

array=[6. 4. 7. 7. 5. 11. 2. 3. 6. 3. 2. 7. 3. 4. 2. 12. 7. 5.]

```
In [27]: ▶ log_fig, ax = plt.subplots()
sns.distplot(np.log(X)).set_title('NaturalLog')
log_fig.savefig('Transformation-NaturalLog' + '.pdf',
                bbox_inches = 'tight', dpi=None, facecolor='w', edgecolor='b',
                orientation='portrait', papertype=None, format=None,
                transparent=True, pad_inches=0.25, frameon=None)
```

C:\Users\gilad\Anaconda3\lib\site-packages\scipy\stats\stats.py:1713: FutureWarning: Using a non-tuple sequence for multidimensional indexing is deprecated; use `arr[tuple(seq)]` instead of `arr[seq]`. In the future this will be interpreted as an array index, `arr[np.array(seq)]`, which will result either in an error or a different result.

```
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```



```
In [ ]: ▶
```

```
In [ ]: ▶
```